

University of Bath Mechanical Engineering Department
Technical Report 13/05:

A new fractal curve

Adrian Bowyer

A.Bowyer@bath.ac.uk

Abstract

The new fractal curve described in this report is deterministic (though it appears at a casual glance to be random), simple to specify, and easy and quick to compute. It will automatically fill any pre-defined bounded and connected region of the plane, is self-avoiding, and never self-intersects. The definition of the curve is given, along with pictures of it in two dimensions generated by a C++ implementation (this software will be freely available on the web on the publication of this work in the reviewed literature). Simulations of the curve are statistically analysed with the result that—under a well-defined and broad range of conditions—its fractal dimension is about 1.5. The curve can use a single line segment inside the bounded region as a starting pattern, or a non-self-intersecting polyline. Several of the curves can be generated in the same region, and will fill it without crossing themselves or each other. The definition extends to any number of dimensions; in K dimensions the curve would become a $K-1$ dimensional hypersurface filling a bounded hypervolume.

Keywords: Peano curve, Hilbert curve, Sierpinski gasket, Koch curve, coastline curve, fractal, space-filling curve.

Introduction

A common computational way of generating an approximation to a random fractal curve (similar to the coastlines discussed by Mandelbrot [7]) is to bisect a straight line segment and to move the half-way point at right angles to the segment by some (pseudo-) random distance. This process is then repeated recursively on the two halves thus created, with the expected distance of the random movement reduced by a fixed factor (say 0.5) at each level of the recursion. Figure 1 from the Boston Center for Polymer Studies [2] shows the result of such a process.




Figure 1: The first five stages of recursion of a random fractal curve from the Boston Center for Polymer Studies [2]. Note the self-intersection near the top.

The method just outlined produces a curve, but that curve is not guaranteed to reside in any particular region of the plane, and nor is it guaranteed never to self-intersect.

Regular space-filling curves such as the Peano [9] and Hilbert [5] curves (Figure 2(a)) never self-intersect, of course, and—if the recursion is infinite—completely cover a known rectangular region of the plane, which is to say that their fractal dimension [8] is 2. There are variations on them which will fill non-rectangular regions as well as rectangular ones (see Figure 2(b)). The region they fill, however, is entirely decided by their elemental shape and the recursion rule.

A self-avoiding curve should also never self-intersect. Peitgen, Jürgens

and Saupe [10] describe a self-avoiding space-filling curve, but this will again only fill a rectangular region. It also prompts the authors to say, “With a plain ‘multiple reduction copy machine’¹ it is not possible to create a self-avoiding and space-filling curve that is aesthetically pleasing.”




Figure 2: (a) The first five stages of recursion of a Hilbert curve [5] from Balayoghan [1]. (b) A regular space-filling curve filling a non-rectangular region.

The recursive curve described in this report is self-avoiding, and is also deterministic like a Peano curve, in that its shape is entirely decided by its initial conditions. But it does not consist of multiple scaled copies of a fixed initial shape (that is, it is not a product of one of Peitgen, Jürgens and Saupe’s multiple reduction copy machines). And I contend that it is aesthetically pleasing².

It also has the useful properties that it will fill any pre-defined connected area of the plane and will never self-intersect³. It is possible to generate it starting from any given non-self-intersecting polyline within the area (including a single line segment). The new curve is also quick and simple to

¹This is an explicatory thought machine (by analogy with a thought experiment) of the authors for producing multiple copies of a shape under transformations.

²Here the reader must make appropriate allowances for authorial bias.

³These two properties were the initial requirements that I had for the curve. They came from a number of engineering applications that were the genesis of it.

compute.

Though this is not investigated here, the definition of the curve extends quite straightforwardly and comprehensively to form (hyper)surfaces in any number of dimensions.

This report is intended to be a descriptive and empirical introduction to the new curve as others may also have applications for it. Should the reader wish to follow this up, I will make C++ software for generating examples of the curve freely available under the GNU Public Licence [4] from my website at <http://staff.bath.ac.uk/ensab/fractal> when the work appears in the reviewed literature. Readers who want a copy before then can e-mail me; I will supply it on the condition that it is not further copied.

The reader wishing to find good books on fractals and space-filling curves is referred to the ones by Hans Sagan [11] and Kenneth Falconer [3].

The new curve

Consider the following method of generating a curve, this time within a bounded area: bisect a straight line segment that lies completely within the area and move the half-way point along the bisector (as in the coastline curve described at the start of the Introduction). But this time move the point *as far as possible as long as it stays at least some exclusion radius away from every other part of the curve and the boundary*. ‘As far as possible’ here means that both halves of the bisector either side of the initial segment need to be considered. The point chosen is the one that lies farthest along the bisector in either direction away from the initial segment.

Now repeat that process recursively on the two halves, reducing the value of the exclusion radius by some factor at each recursive step, again as before. In general the exclusion is for all parts of the curve except the split-line-segment’s two immediately-joining before-and-after line segments, as its distance to them must always be zero.

The result is a curve such as the one shown in Figure 3, which was generated by the C++ program mentioned above.




Figure 3: An example of the new curve. The closed area was the unit square, and the starting line segment was its diagonal from the origin. The initial exclusion radius was 0.1, and this was reduced by a factor of 0.8 at each level of recursion. The recursion was run 10 times.

Characteristics of the new curve

Figure 4 shows the original square diagonal and the first five recursions in the generation of the curve in Figure 3. Clearly the order in which the recursion is done will affect the resulting curve. This example used breadth-first recursion, and the curve was alternately scanned from bottom left to top right, and then back again. If the example in Figure 3 is computed breadth-first, but always scanned from bottom left to top right then the result is as in Figure 5. Note how this has led to large line segments near the origin and smaller ones near the top right. More on this below. If the recursion were to be done depth-first this would generate a much greater bias, as, after the first division, the first half of the original line would be fully divided before the second half was considered as anything more than an obstruction. The choice of recursion sequence is one of the freedoms available in the definition of the curve.

The first division in Figure 4 obviously involves a degenerate choice between two possibilities in that—because of symmetry—the split-point could equally have been moved to the bottom-right corner of the square. Note that, even though the result of that first division is almost as symmetrical,




Figure 4: The original segment and the first five recursions of the curve in Figure 3. The recursion was done breadth-first, and the polyline being generated was alternately scanned from bottom left to top right and then back again. The discs are the exclusion radii at each recursion level.

the second division is completely determined by the order in which the segments of the polyline were visited. Other than such degeneracies (which in general will almost-surely never happen), the curve is completely determined by the starting conditions, the recursion sequence, and the factor by which the exclusion radius is reduced at each step.

Two other freedoms available in the definition of the curve are the choice of that initial exclusion radius and its reduction factor. Figure 6 (a) shows the curve resulting from the same conditions as Figure 3, but with a reduction factor of 0.7 instead of 0.8.

The ‘gaps’ in the pattern (more pronounced in Figure 6 than in Figure 3) are filled by subsequent recursions as illustrated by Figure 6 (b), where the line segment labelled A in the left-hand figure has split with the resulting point moving across the triangle-shaped void to near the opposing vertex. This gives rise to an interesting phenomenon: the stability of near-




Figure 5: The same conditions as Figure 3, except that the curve was always scanned from bottom left to top right.




Figure 6: (a) The same conditions as Figure 3, except that the exclusion radius reduction factor was 0.7 instead of 0.8. (b) The next level of recursion down, in which the line segment labelled A has been split and the split point moved to near the opposite vertex of the triangle-shaped void.

equilateral-triangle-shaped voids in patterns with low values of the exclusion-radius reduction-factor. When such a shape arises, the stability happens because one of its edges gets split and the split point moves to near the opposite vertex, roughly preserving the shape. This is illustrated in Figure 7, which

shows the first few stages of recursion on an actual equilateral triangle inside a slightly larger equilateral boundary.




Figure 7: The stability of an equilateral triangle: the starting configuration, the first four recursions, and the tenth recursion.

High values of the reduction factor avoid this phenomenon for intuitively obvious reasons: the segments are being split and becoming shorter and shorter, but the exclusion radius is not falling so fast with the recursion. The curve therefore is becoming more flexible and string-like (as opposed to being like a chain of hinged rods), and its freedom of movement is constrained by the comparatively high exclusion radius values—it threads its way between other parts of itself near to the median boundaries defined by the exclusion discs.

Simulation and statistics

If the curve were a true fractal, one would expect its length to increase with the depth of recursion as a power law, which is to say that one would expect

$$\ln(L) = d \ln\left(\frac{2^N}{L_0}\right)$$

where L is a measured length of the curve, L_0 is the length of the original line that was split to make the curve, N is the depth of recursion ($N = 0$ at $L = L_0$), and d is the power-law exponent. The value of $\frac{2^N}{L_0}$ is the number of times the first segment has been cut divided by its original length. In other words, it is an inverse resolution-length for the total length measurement. The fractal dimension of the curve, D , would be given by $D = d + 1$.




Figure 8: The length of the curve as a power law. The top characteristics are virtually indistinguishable and correspond to the conditions in Figure 3 with reduction factors of the exclusion radius of 0.3, 0.5, 0.6 and 0.7. The remaining characteristics have the reduction factors indicated at their ends.

Figure 8 shows plots of $\ln(L)$ against $\ln(\frac{2^N}{L_0})$ for a range of radius reduction factors with the conditions otherwise as in Figure 3. The data points have been omitted for clarity, but are equally-spaced along the abscissa for each characteristic and correspond to each value of N from 0 to 15. L_0 , in this case, is $\sqrt{2}$.

For reduction factors of 0.7 and below, the characteristics are straight and virtually indistinguishable.

It is clear that the line for a reduction factor of 0.95 is not straight, though it is not clear whether its asymptote would be finite. A finite asymptote would imply that, even after infinite recursion, the corresponding curve would still have a finite length. A non-straight characteristic with an infinite value after infinite recursion would imply a curve of infinite length, but not a curve that was a scale-free self-similar fractal. For scale-free self-similarity the characteristic line has to be straight, like those for reduction factors of 0.7 and below.

If each characteristic for reduction factors of 0.7 and below is subjected to linear regression the average of the gradients, d , of the resulting straight lines is 0.496 with a standard error of 0.002. This means that for those four characteristics the mean fractal dimension is 1.496, which—with the necessary caveats about small sample size—tempts me to the intriguing conjecture that the true value may be exactly 1.5.

Extending the idea




Figure 9: A curve that started as a square of half the side-length of the bounding box placed at the box’s centre. The flood fill is to make the curve-generating method’s preservation of the interior and the exterior clearer.

Instead of starting with a single line segment, one can start with a pre-defined polyline (or polygon, as with the triangle in Figure 7). Figure 9 shows the result of starting with a square of side-length 0.5 in the middle of the unit bounding box. The resulting closed curve has been flood-filled to make the curve-generating-method's preservation of the interior and the exterior clearer.




Figure 10: The curve in an L-shaped bounding region; the starting curve was also L-shaped along the centre lines of the two limbs of the bounding region.

The bounding region can be any shape. Figure 10 shows a curve generated in an L-shaped bounding area. This leads immediately to the idea of running several intertwining curves at once. Figure 11 shows the starting curves and the resulting curves from such an experiment.

Conclusions

The curve that I needed would:

1. Fill any pre-defined closed connected area;
2. Be unbroken;
3. Not self-intersect;




Figure 11: The starting configuration, and the result of running the two starting polylines against each other.

4. Be possible to generate starting from any given non-self-intersecting polyline within the area (including, obviously, any single line segment);
5. Extend straightforwardly and comprehensively to any number of dimensions (especially 3); and
6. Be reasonably quick and simple to compute for finite recursions.

The curve presented in this report fulfils all those requirements.

Extending the curve in two-dimensions up to a surface in three-dimensions is conceptually simple: the closed area would become a closed volume (usually, though not necessarily, bounded by a polyhedron) and the starting line segment (or segments) would become triangle(s) within it. There are a number of ways that a triangle can be split into other triangles. The simplest, perhaps, is to bisect each edge and to join those points to form four new triangles. (This scheme is commonly used to generate random fractal landscapes, as described by Hughes [6].) The movement of the splitting-points would be slightly more complicated than in two dimensions: for splitting-points on edges that were interior to the surface being created the points would be moved as far as possible along a line that was perpendicular to the edge and that bisected the dihedral angle between the two triangles that shared the edge; for splitting points on edges at the boundary of the surface being created the points would be moved as far as possible along a line that

was perpendicular to the edge and that lay in the plane of the triangle owning the edge. In both cases the points would be (as in two dimensions) moved so that they and the triangles they defined came no closer than the exclusion radius to anything else. In K dimensions the triangles would become $K - 1$ dimensional simplexes and the boundary would usually (though again, not necessarily) be a polytope.

The next piece of work that I will do on this will be to write software for constructing self-avoiding surfaces in three dimensions, as just described.

I would be most interested to hear of possible applications for the new curve. Also, as I intend to work on applications of it rather than analysis of it myself, I would be interested to hear of analytical results on its characteristics. My e-mail address is given above.

Acknowledgement

I would like to thank John Woodwark for his comments on the first draft of this report.

References

- [1] V. B. Balayoghan (Department of Computer Sciences, The University of Texas at Austin): applet for generating various regular Peano curves at <http://www.cs.utexas.edu/users/vbb/misc/sfc/Oindex.html>
- [2] Boston University Center for Polymer Studies: Coastline applet at <http://polymer.bu.edu/java/java/coastline/coastline.html>
- [3] Kenneth Falconer: *Fractal Geometry: Mathematical Foundations and Applications*, Wiley, ISBN 0 471 92287 0, 1990.
- [4] Free Software Foundation Inc. (59 Temple Place - Suite 330, Boston, MA 02111-1307, USA): *GNU General Public License* at <http://www.gnu.org/copyleft/gpl.html>
- [5] David Hilbert: *Über die stetige Abbildung einer Linie auf ein Flächenstück*, Math. Ann., **38**, pp. 459-460, 1891.

- [6] Merlin Hughes: *3D Graphic Java: Render fractal landscapes*, at <http://www.javaworld.com/javaworld/jw-08-1998/jw-08-step-p.html>
- [7] Benoit B. Mandelbrot: *How Long is the Coast of Great Britain: Statistical Self Similarity and Fractional Dimension*, Science, **156**, 636-638, 1967.
- [8] Benoit B. Mandelbrot: *Fractals: Form, Chance, and Dimension*, Freeman, ISBN 0716704730, 1977.
- [9] Giuseppe Peano: *Sur une courbe, qui remplit toute une aire plane*, Math. Ann. **36**, pp. 157-160, 1890.
- [10] Heinz-Otto Peitgen, Hartmut Jürgens, & Dietmar Saupe: *Chaos and Fractals – New Frontiers of Science*, Springer-Verlag, ISBN 0-387-97903-4, 1992.
- [11] Hans Sagan: *Space-Filling Curves*, Springer-Verlag, New York, ISBN 0-387-94265-3, 1991.